

# Implementing an Efficient Shuffle Operator for Streaming Database Systems

**Bachelor Thesis**

**Author:** Jonas Ladner

**Supervisor:** Maximilian Rieger, M.Sc.

**Examiner:** Prof. Dr. Thomas Neumann



Garching, 11.03.2025

Technical University of Munich

# Problem Setting

**Key Contribution:** Efficient, multithreaded shuffle operator implementations.

## Shuffle-Simulation Process:

1. **Tuple Generation:** Randomly generated tuples with 32-bit keys and optional data fields.
2. **Data Shuffle:** Tuples stored in partition buckets using slotted pages.
3. **Storing on Slotted Pages:** Thread-local vs. shared (locking/lock-free) write-out strategies.

# Naive approach: *OnDemand*

- **OnDemand:** Tuples are directly written to the partition buckets.
- **Problems:**
  - High contention on partition buckets.

# Optimized approach: *Smb*

- **Software Managed Buffers (SMBs):**
  - Cacheline-sized, thread-local buffers for each partition.
  - Flush when buffer is full.
- **Problems:**
  - High contention on partition buckets.

# Histogram-based approach: *Radix*

- **Histogram:**
  - Thread-local histograms for each partition.
  - Flush when histogram is full.
- **Problems:**
  - High contention on histograms.

# Histogram-based approaches: *Radix* and *Hybrid*

- **Histogram:**
  - Thread-local histograms for each partition.
  - Flush when histogram is full.
- **Problems:**
  - High contention on histograms.

# Reducing contention: *CmpProcessingUnits*

# Avoiding contention: *LocalPagesAndMerge*



# Evaluation




Figure: Benchmark Plots for Tuple of 16B with 2 and 8 Partitions

# Evaluation




Figure: Benchmark Plots for Tuple of 4B with 32 and 1024 Partitions

# Evaluation




Figure: Benchmark Plots for Tuple of 16B with 32 and 1024 Partitions

# Evaluation

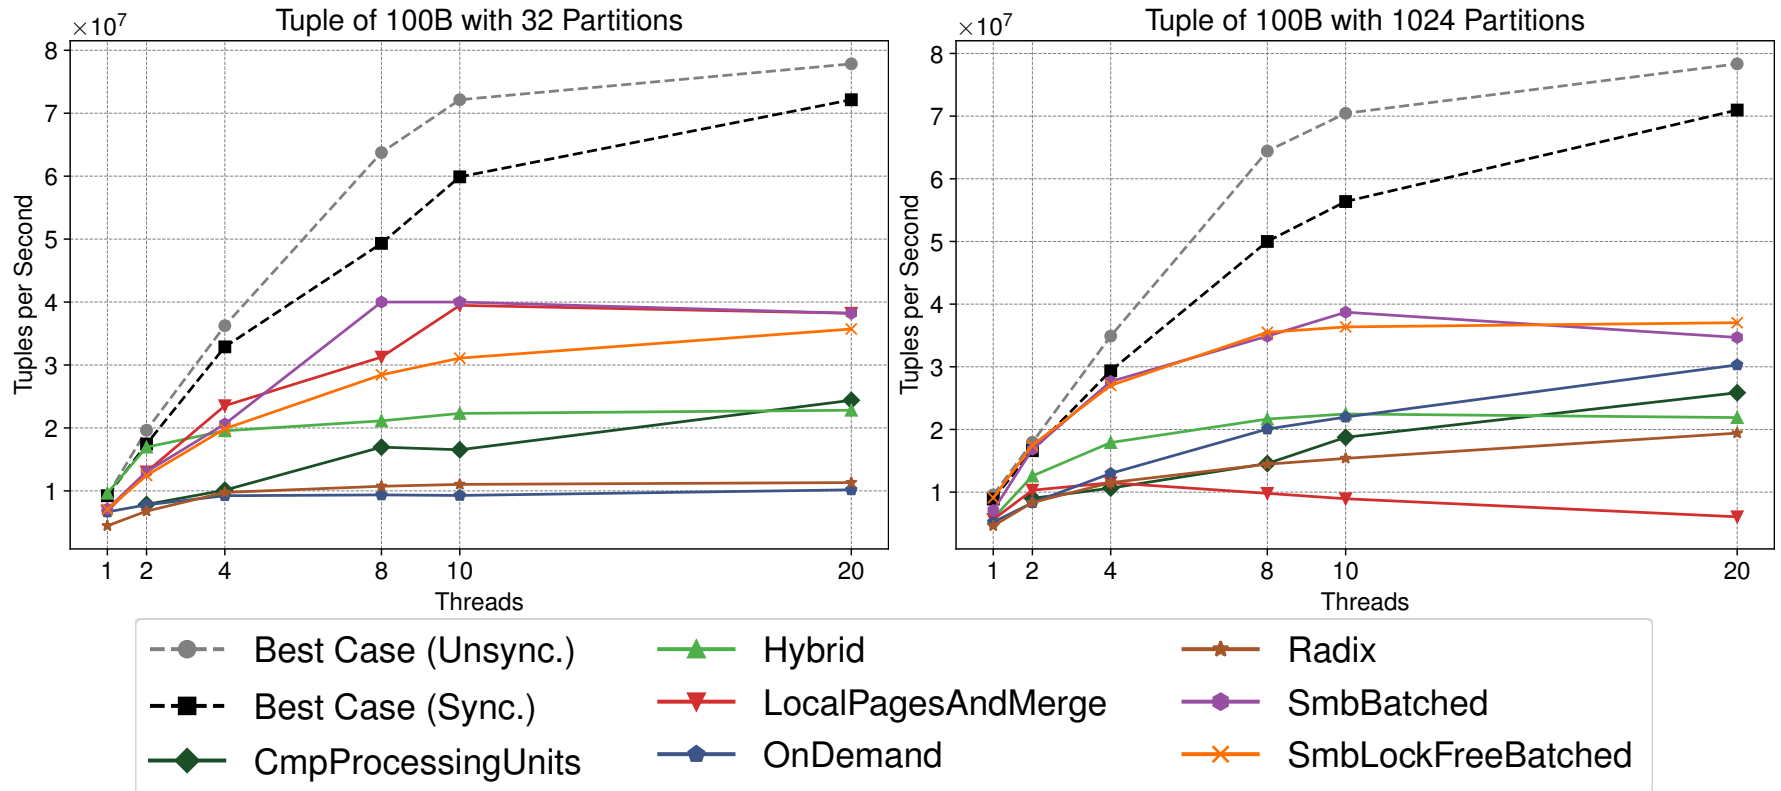


Figure: Benchmark Plots for Tuple of 100B with 32 and 1024 Partitions

# Peak Heap Memory




Figure: Peak Heap Usage when using 32 Partitions, 40 Threads and 67.2 Mio. 16B Tuples (1 GiB)

# Comparison with Apache Flink




Figure: Tuples per Second Comparison when using 1024 Partitions

# Conclusion and Future Work

## Conclusion:

- Thread-local slotted pages are optimal for low partition counts.
- SMB-based methods scale best for high partition counts.

## Future Work:

- Further reducing contention on machines with 20+ cores.
- Evaluate the implementations in a real-world streaming system.