



SCHOOL OF COMPUTATION,  
INFORMATION AND TECHNOLOGY —  
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

# **Implementing an Efficient Shuffle Operator for Streaming Database Systems**

Jonas Ladner





SCHOOL OF COMPUTATION,  
INFORMATION AND TECHNOLOGY —  
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Implementing an Efficient Shuffle Operator  
for Streaming Database Systems**

**Implementierung eines effizienten Shuffle-  
Operators für Streaming-Datenbanksysteme**

|                  |                          |
|------------------|--------------------------|
| Author:          | Jonas Ladner             |
| Examiner:        | Prof. Dr. Thomas Neumann |
| Supervisor:      | Maximilian Rieger M.Sc.  |
| Submission Date: | 17.02.2025               |

I confirm that this bachelor's thesis is my own work and I have documented all sources and material used.

Munich, 17.02.2025

Jonas Ladner

## **Acknowledgments**

# Abstract

Modern streaming database systems rely on efficient data partitioning to achieve scalability and high performance across processing nodes. Partitioned data shuffling is a crucial operation, as it is used to prepare and distribute data for further processing on distributed systems.

This thesis purposes and evaluates various partitioning implementations by simulating real-world usage of the shuffle operator. The implementations process incoming tuple batches and partition them into output buckets, which are based on slotted pages and can be passed to subsequent operators. The evaluation of the implementations is based on their performance, scalability and memory consumption.

The results demonstrate that using a lock- and Software Managed Buffer (SMB)-based approach yields the best overall performance, offering both efficiency and ease of implementation. Notably, the proposed locking mechanism minimizes the duration of holding a lock, ensuring minimal contention and contributing to the approach's superior performance.

TODO: quantify results

"As a result we implement a ... that is ?X faster than ..."

At the end: "Our approach uses software managed buffers, locking, ..."

# Contents

|   |            |
|---|------------|
| <b>Acknowledgments</b>  | <b>iii</b> |
| <b>Abstract</b>   | <b>iv</b>  |
| <b>1 Introduction</b>   | <b>1</b>   |
| 1.1 Motivation . . . . .  | 1          |
| 1.2 Streaming processing engines . . . . .                      | 1          |
| 1.3 Shuffle operator . . . . .                                  | 1          |
| 1.4 Slotted pages . . . . .                                     | 2          |
| 1.5 Problem setting . . . . .                                   | 2          |
| <b>2 Related work</b>   | <b>4</b>   |
| 2.1 Radix Partitioning . . . . .                                | 4          |
| 2.2 Partitioned Joins . . . . .                                 | 4          |
| 2.3 Software Managed Buffers . . . . .                          | 4          |
| <b>3 Implementations</b>  | <b>5</b>   |
| 3.1 Partitioning . . . . .                                      | 5          |
| 3.2 Slotted Pages . . . . .                                     | 5          |
| 3.3 Slotted Page Managers . . . . .                             | 5          |
| 3.3.1 Lock-based Page Manager . . . . .                         | 6          |
| 3.3.2 Lock-free Page Manager . . . . .                          | 7          |
| 3.3.3 Histogram-based Page Managers . . . . .                   | 8          |
| 3.3.4 Thread-Local Pages and Merge-based Page Manager . . . . . | 10         |
| 3.3.5 Implementation-independent Optimizations . . . . .        | 10         |
| 3.4 On-Demand Partitioning . . . . .                            | 11         |
| 3.4.1 Overview . . . . .  | 11         |
| 3.5 Software Managed Buffers-based Partitioning . . . . .       | 12         |
| 3.5.1 Overview . . . . .  | 12         |
| 3.5.2 Lock-based Implementations . . . . .                      | 12         |
| 3.5.3 Lock-free Implementations . . . . .                       | 12         |
| 3.6 Histogram-based Partitioning . . . . .                      | 12         |
| 3.6.1 Overview . . . . .  | 12         |

|          |  |           |
|----------|--|-----------|
| 3.6.2    | Radix Partitioning . . . . .                                     | 12        |
| 3.6.3    | Ad-hoc Radix ("Hybrid") Partitioning . . . . .                   | 12        |
| 3.7      | Thread-Local Pages and Merge-based Partitioning . . . . .        | 12        |
| 3.7.1    | Overview . . . . .   | 12        |
| 3.8      | Collaborative Morsel Processing . . . . .                        | 12        |
| 3.8.1    | Overview . . . . .   | 12        |
| 3.8.2    | Collaborative Morsel Processing using Software Managed Buffers   | 12        |
| 3.8.3    | Collaborative Morsel Processing using Processing Units . . . . . | 12        |
| 3.9      | Complexity Analysis . . . . .                                    | 12        |
| 3.9.1    | Time Complexity . . . . .  | 12        |
| 3.9.2    | Space Complexity . . . . .                                       | 12        |
| <b>4</b> | <b>Evaluation</b>  | <b>13</b> |
| 4.1      | Experimental Setup . . . . .                                     | 13        |
| 4.1.1    | Hardware . . . . .   | 13        |
| 4.1.2    | Software . . . . .   | 13        |
| 4.2      | Tuple Generation . . . . .                                       | 13        |
| 4.3      | Tuple Write Benchmark . . . . .                                  | 13        |
| 4.4      | Shuffle Benchmark . . . . .                                      | 13        |
| 4.4.1    | Memory Consumption . . . . .                                     | 13        |
| 4.4.2    | Performance . . . . .  | 13        |
| 4.5      | Comparison with Stream Processing Systems . . . . .              | 13        |
| <b>5</b> | <b>Conclusion</b>  | <b>14</b> |
| 5.1      | Conclusion . . . . .   | 14        |
| 5.2      | Future Work . . . . .  | 14        |
|          | <b>Abbreviations</b>   | <b>15</b> |
|          | <b>List of Figures</b>   | <b>16</b> |
|          | <b>List of Tables</b>  | <b>17</b> |
|          | <b>List of Algorithms</b>  | <b>18</b> |
|          | <b>Bibliography</b>  | <b>19</b> |

# 1 Introduction

## 1.1 Motivation

Growing demand for real-time data analysis and increasing data volume create significant challenges [1–3]. Distributed and parallel systems like stream processing engines address these issues by distributing tasks across worker nodes [1, 4]. Data shuffling prepares and distributes tuples among worker nodes [5]. As a core streaming component, the shuffle operator must achieve high throughput and low latency to support efficient downstream processing. This thesis addresses these challenges by focusing on the most efficient implementations of the shuffle operator.

## 1.2 Streaming processing engines

Streaming processing engines are designed to process data as soon as it arrives rather than relying on traditional pre-computed information and index structures [1]. Their core operations include partitioning and distributing incoming traffic across worker nodes. The distribution process is frequently based on a partitioning function. These partitions can then improve the performance of further operators by ensuring the data locality of interdependent tuples [3, 6]. Data locality within the worker node is crucial for maintaining performance and scalability in large-scale deployments.

## 1.3 Shuffle operator

The shuffle operator provides a partitioned distribution of tuples, enabling downstream operators to leverage the data locality of partitioned data blocks. For instance, the throughput of the join operator can significantly be improved when tuples assigned to the same hash bucket are shuffled to the same worker node [3]. Implementing a shuffle operator involves addressing challenges like memory consumption, latency, and scalability. This thesis proposes and evaluates different implementations of the shuffle operator, focusing on their efficiency and performance.



## 1.4 Slotted pages

Slotted pages are a common way to store variable-size tuples within fixed size memory blocks [7]. These fixed size memory blocks can then be either stored on disk or easily be sent to worker nodes.

Typically, the pages consist of three sections: metadata, slots and a variable-size data section. The metadata area contains information like an identifier for the page, what fields the tuples have and the amount of tuples on this page. This fixed-size metadata section is then followed by the slot section. A single slot contains the fixed-size properties, the variable-size length and its start offset in the variable-size data section. In contrast to the previous two sections, the variable-size data section grows from the end of the page towards the slot section of the page.

## 1.5 Problem setting

The key contribution of this thesis is the creation and evaluation of the most efficient, multithreaded implementations of the shuffle operator. In order to simulate the real-world usage of the shuffle operator in a streaming system, the following three-step shuffle-simulation is proposed:

1. Tuple generation: The tuples are generated in a batched manner using a pseudo-random generator. Each implementation requests the ad-hoc generation of tuples, which contain a 32-bit key field and an optional variable-size data field.
2. Data shuffle: The requested, random-generated tuples are then processed using the different implementations and stored within partition buckets. Each partition bucket consists of slotted pages, where the tuple of this partition are stored.
3. Storing tuples on slotted pages: The implementations range from thread-local to shared slotted-page write-out strategies. The implementations using a shared write-out policy, can then be further categorized into locking and lock-free approaches.

This simulation is close to the real world usage of the shuffle operator, as streaming systems work on incoming tuple batches that are not materialized like in traditional relational databases.

Further, we use slotted pages as a communication format between nodes. This is an efficient and widely used format to transport tuples within fixed size data chunks. In contrast, sending each tuple as soon as it is processed, creates significant network overhead and makes downstream processing less efficient.

While the implementations are optimized for the simulated process above, the underlying algorithms can be transferred to any streaming system, that forwards data using slotted pages.

## **2 Related work**

### **2.1 Radix Partitioning**

### **2.2 Partitioned Joins**

### **2.3 Software Managed Buffers**

## 3 Implementations

This section explains the implementations and how to implement them efficiently. As the shuffle operator-simulation is based on fixed-size tuples, the explanations of the following implementation, are based on fixed-size tuples as well.

### 3.1 Partitioning

We generate our tuples using an pseudo-random number generator. To distribute the tuples into partitions, we use the following function:

$$f(t) = t.\text{key} \% \text{partitions} \quad (3.1)$$

As the modulo calculation is quite expensive, it we use the following partitioning function, when the count of partitions is a power of two:

$$f(t) = t.\text{key} \& (\text{partitions} - 1) \quad (3.2)$$

### 3.2 Slotted Pages

Slotted pages store their information on fixed size memory blocks, that are split up in three sections: a fixed-size header, slots and a data section. We are using slotted pages with a total size of 5 MiB per page.

In our implementation, we only store the tuple count in the header. To construct our shuffle-simulation close to the real world usage, we split up each tuple. In the slot information, we store the 4-Byte key together with data offset and length information. The remainder of the tuple is then stored in the data section at the end of each page.

As we are using fixed-size tuples in our simulation, we only need to have the tuple slot index to be able to store the tuple on the page. In constrast, when dealing with variable size tuples, slot index and data offset are needed to store a tuple.

### 3.3 Slotted Page Managers

As some implementation share the same tuple write-out strategy, we propose the used write-out strategies here and reference them in the following explanations of the



Figure 3.1: Slotted Page grow visualization

concrete implementations.

To further simplify the implementations, we initialize each partition with an empty slotted page. This significantly reduces the complexity of the page manager implementations.

### 3.3.1 Lock-based Page Manager

For each partition, we use a single lock and a vector for storing the slotted pages. As

---

#### Algorithm 1: Lock-based Page Manager insert\_tuple Algorithm

---

**input** : tuple: The tuple to be inserted, partition: The target partition index  
**output**: Tuple inserted into the appropriate slotted page of the specified partition.

```

1 function insert_tuple(tuple, partition)
2   Acquire lock on partition_locks[partition]
3   if pages[partition].back().add_tuple(tuple) then
4     // Tuple added successfully
5   else
6     add_page(partition)
7     pages[partition].back().add_tuple(tuple)
8   end
9   Release lock

```

---

can be seen in Algorithm 1, the lock-based insertion process is straightforward. The insertion on a given slotted page, can only fail if the page is full. This can easily be checked by reading the tuple count in the metadata section of the slotted page. If the current page is full, we just allocate and append a new slotted page to the page vector of this partition.

### Tuple insertion in batches

Similarly, we can further optimize the write-out by using tuple-batches. In Algorithm 2,

---

**Algorithm 2:** Lock-based Page Manager insert\_tuple\_batch Algorithm

---

```

input : tuples: The tuple-batch to be inserted
        partition: The target partition index
output: Tuples inserted into one or more slotted pages of the specified partition.
1 function insert_tuple_batch(tuples, partition)
2   Acquire lock on partition_locks[partition]
3   for tuple : tuples do
4     if pages[partition].back().add_tuple(tuple) then
5       | // Tuple added successfully
6     else
7       | add_page(partition)
8       | pages[partition].back().add_tuple(tuple)
9     end
10 end
11 Release lock

```

---

we reuse the tuple insertion logic from Algorithm 1 but acquire the partition lock only once for the entire insertion process. Since acquiring and releasing the lock is expensive, this optimization significantly improves performance in multi-threaded scenarios.

### 3.3.2 Lock-free Page Manager

As holding a lock of a partition denies a second thread to also write out tuples, we propose a lock-free implementation. In comparison to the lock-based variant, we now have to store our slotted pages in a pointer-stable data structure. This is necessary to ensure threads can work simultaneous, while a thread adds a new slotted page. Furthermore, we have to edit the slotted page metadata using compare-and-exchange operations to avoid losing writes from other threads.

In order to gather the information, where we can write a tuple, we increment the tuple count using compare-and-exchange. This index then acts as our location, where our tuple is placed on the page. In Algorithm 3, we also add an condition to stop attempting to further increase the count of tuples on the page, if the maximum is reached. This ensures that threads move to the next allocated page. Given the index, where the tuple is placed, we also append a pointer of the start of the page and the page size. This ensures we can write the page without requiring any further information.

---

**Algorithm 3:** Lock-free Slotted Page increment\_and\_fetch\_opt\_write\_info Algorithm

---

```

1 function increment_and_fetch_opt_write_info()
2   current_tuple_count = header->tuple_count.load();
3   while !header->tuple_count.compare_exchange_strong(current_tuple_count,
      current_tuple_count + 1) do
4     if current_tuple_count >= get_max_tuples(page_size) then
5       |   return std::nullopt
6     end
7   end
8   return {page_data.get(), page_size, current_tuple_count}

```

---



---

**Algorithm 4:** Lock-free Page Manager insert\_tuple Algorithm

---

```

input : tuple: The tuple to be inserted, partition: The target partition index
output: Tuple inserted into the appropriate slotted page of the specified partition.
1 function insert_tuple(tuple, partition)
2   wi = current_page[partition].load()->increment_and_fetch_opt_write_info()
3   while wi == std::nullopt do
4     |   wi = current_page[partition].load()->increment_and_fetch_opt_write_info()
5   end
6   if wi.tuple_index == LockFreeSlottedPage::get_max_tuples() - 1 then
7     |   add_page(partition)
8   end
9   LockFreeSlottedPage::add_tuple_using_index(wi, tuple)

```

---

Using the Algorithm 3, we retrieve the information to write the tuple in Algorithm 4. We read from an atomically-stored pointer to our current page, until we are assigned an index on a slotted page. If we are writing the last tuple on the page, we add a new page to this partition. This is necessary to create a unique condition, when a new page has to be allocated. With that write information at hand, we can write the tuple onto the slotted page.

### 3.3.3 Histogram-based Page Managers

The following page managers are based on histograms, which typically store tuple count and total tuple size per partition. Similar to Section 2.1, this information can then be used to assign memory areas on slotted pages.

### **Radix Page Manager**

This page manager applies the concept of radix-partitioning on slotted pages. It uses an three step process:

1. Histogram retrieval: Each thread reads its assigned materialized tuple chunk and builds up an histogram. As we are using fixed-size tuples, it only stores the tuple count per partition. This histogram is then forwarded to the Radix Page Manager, which collects the histogram of each thread and then moves on to step 2.
2. Page allocation: With all histograms at hand, the page manager sums up each histogram into a global histogram. This global histogram stores how many tuples each partition has to store. With that information we can allocate required pages for each partition. When all pages are ready to be used, step 3 begins.
3. Assignment of slotted page sub-chunks: Each thread uses its local histogram to request storage locations for its tuples. The page manager uses the pre-allocated pages to assign each thread one or more memory chunks. These memory chunks can then be used exclusively by the thread to store each tuples. Only the tuple count in the metadata has to be atomically updated once, to signal that this thread has finished its work on this page. Otherwise, this write-out process does not rely on any synchronisation after the memory chunks.

After these three steps are done, the page manager is finished. Each page can be send to its receiver, once the expected tuple count is reached. This can be done by the last thread to increment the tuple count.

### **Ad-hoc Radix ("Hybrid") Page Manager**

The idea of the previous Radix Page Manager can be used to construct an approach, where each thread can hand in its histogram and receive memory chunks on slotted pages independently from other threads. This page manager merges the three steps of the Radix Page Manager into one.

A thread hands in its histogram and request the memory chunks, where the tuples can be written. The page manager reads the histogram and processes each partition individually. For each partition, where a tuple has to be stored, the lock of the partition is aquired. If there is any space left on a partition, it is used up and assigned to this thread. When the current page is full, a new page is allocated and the page manager continues the assignment process. These exclusive memory locations are then used by the thread to store its tuples.

In comparision to the Radix Page Manager, this page manager does not pre-allocate all necessary pages. Instead it allocates the pages when needed. Furthermore, a thread



can already receive its memory chunks, while other threads are still constructing their histograms. This allows this implementation to avoid materialization of all tuples.

### 3.3.4 Thread-Local Pages and Merge-based Page Manager

A further step into reducing the necessity of synchronization are thread-local slotted pages. This slotted page management scheme can be split into two phases:

1. Thread-local write-out: These slotted pages are exclusively used by the owning thread and after the tuple processing is done, each thread hands in their pages. When the page manager has received all pages of each thread, the merging phase is started.
2. Page merging: The page manager splits up all partitions onto the available threads. Each thread is then responsible for merging the slotted pages for each partition in the assigned partition range. To minimize tuple movement, the slotted pages are sorted decreasingly by the tuple count. Then the pages with fewer tuples are merged into the fuller pages. It can be the case that the last page cannot be fully merged into any other slotted page. Then this page has to be reordered so that the slots and data section start at the section beginnings. This is necessary to avoid having a gap at the beginning of the slot or data section.

After the merging phase, all tuples are stored on the least possible amount of slotted pages. During the thread-local processing, it is likely that more slotted pages than necessary are created. This can lead to a significant higher memory consumption than the previous approaches.

### 3.3.5 Implementation-independent Optimizations

We use the following optimizations to speed up our simulation of the shuffle operator.

#### **Padded atomics and locks**

All implementations, that use an array of partition locks, are affected by false sharing. False sharing is caused when a cache line stores two independent values and one Central Processing Unit (CPU) core is modifying the first value. Then another CPU core wants to access the second value, which causes a cache miss.

This performance degradation can be avoided by storing each individual partition lock aligned to the L1-cache line boundary. This significantly reduces the amount of L1-cache misses, as the partition locks are frequently accessed.

### **Minimal page-locking**

When holding a partition lock, the tuple write-out onto a slotted page is an expensive operation. To reduce the lock-duration, we gather the necessary write-out information and prepare the next tuple insertion, and releasing the lock before the actual tuple write-out.

### **Two-step buffered slotted page write out**

When writing out a batch of tuples onto a slotted page, we can process each tuple individually. First, we construct the slot for the tuple, and then we store the variable-size data in the data section at the end of page. Linux on an x86-64 machine, typically uses 4 KiB memory pages. As we use 5 MiB slotted pages, it is expected that the slot and data are on different memory pages.

To reduce the amount simultaneous accessed pages, we store our batches in a slot phase and a data phase. As we are now writing to either the slot or data section, we switch between memory pages less often. This is more friendly to the CPU cache and the Translation Lookaside Buffer (TLB).

## **3.4 On-Demand Partitioning**

On-demand Partitioning is the simplest algorithm to implement the shuffle operator.

### **3.4.1 Overview**

As soon as this implementation receives a batch of tuples, it processes each tuple individually. First, the hash-function is applied onto the tuple to gather information in which partition this tuple belongs. With this information, we only need to write out the tuple into the corresponding partition. The partition is now locked, and we check, if there is enough space left on the newest slotted page of this partition.

## **3.5 Software Managed Buffers-based Partitioning**

### **3.5.1 Overview**

### **3.5.2 Lock-based Implementations**

### **3.5.3 Lock-free Implementations**

## **3.6 Histogram-based Partitioning**

### **3.6.1 Overview**

### **3.6.2 Radix Partitioning**

### **3.6.3 Ad-hoc Radix ("Hybrid") Partitioning**

## **3.7 Thread-Local Pages and Merge-based Partitioning**

### **3.7.1 Overview**

## **3.8 Collaborative Morsel Processing**

### **3.8.1 Overview**

### **3.8.2 Collaborative Morsel Processing using Software Managed Buffers**

### **3.8.3 Collaborative Morsel Processing using Processing Units**

## **3.9 Complexity Analysis**

### **3.9.1 Time Complexity**

**Tuple Access Count**

### **3.9.2 Space Complexity**

**Memory Consumption**

## **4 Evaluation**

### **4.1 Experimental Setup**

#### **4.1.1 Hardware**

#### **4.1.2 Software**

### **4.2 Tuple Generation**

### **4.3 Tuple Write Benchmark**

### **4.4 Shuffle Benchmark**

#### **4.4.1 Memory Consumption**

#### **4.4.2 Performance**

### **4.5 Comparison with Stream Processing Systems**

## **5 Conclusion**

### **5.1 Conclusion**

### **5.2 Future Work**

# Abbreviations

**CPU** Central Processing Unit

**TLB** Translation Lookaside Buffer

**SMB** Software Managed Buffer

# List of Figures

|     |   |   |
|-----|---|---|
| 3.1 | Slotted Page grow visualization . . . . . | 6 |
|-----|---|---|

## List of Tables



## List of Algorithms

|   |   |   |
|---|---|---|
| 1 | Lock-based Page Manager insert_tuple Algorithm . . . . .              | 6 |
| 2 | Lock-based Page Manager insert_tuple_batch Algorithm . . . . .        | 7 |
| 3 | Lock-free Slotted Page increment_and_fetch_opt_write_info Algorithm . | 8 |
| 4 | Lock-free Page Manager insert_tuple Algorithm . . . . .               | 8 |

# Bibliography

- [1] F. Gürcan and M. Berigel. “Real-Time Processing of Big Data Streams: Lifecycle, Tools, Tasks, and Challenges.” In: *2018 2nd International Symposium on Multi-disciplinary Studies and Innovative Technologies (ISMSIT)*. Oct. 2018, pp. 1–6. DOI: 10.1109/ISMSIT.2018.8567061.
- [2] E. Zamanian, C. Binnig, and A. Salama. “Locality-aware Partitioning in Parallel Database Systems.” In: *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*. SIGMOD ’15. Melbourne, Victoria, Australia: Association for Computing Machinery, 2015, pp. 17–30. ISBN: 9781450327589. DOI: 10.1145/2723372.2723718.
- [3] S. Chu, M. Balazinska, and D. Suciu. “From Theory to Practice: Efficient Join Query Evaluation in a Parallel Database System.” In: *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*. SIGMOD ’15. Melbourne, Victoria, Australia: Association for Computing Machinery, 2015, pp. 63–78. ISBN: 9781450327589. DOI: 10.1145/2723372.2750545.
- [4] G. Andrade, D. Griebler, R. Santos, and L. G. Fernandes. “A parallel programming assessment for stream processing applications on multi-core systems.” In: *Computer Standards & Interfaces* 84 (2023), p. 103691. ISSN: 0920-5489. DOI: <https://doi.org/10.1016/j.csi.2022.103691>.
- [5] F. Liu, L. Yin, and S. Blanas. “Design and Evaluation of an RDMA-aware Data Shuffling Operator for Parallel Database Systems.” In: *ACM Trans. Database Syst.* 44.4 (Dec. 2019). ISSN: 0362-5915. DOI: 10.1145/3360900.
- [6] W. Rödiger, T. Mühlbauer, P. Unterbrunner, A. Reiser, A. Kemper, and T. Neumann. “Locality-sensitive operators for parallel main-memory database clusters.” In: *IEEE 30th International Conference on Data Engineering, Chicago, ICDE 2014, IL, USA, March 31 - April 4, 2014*. Ed. by I. F. Cruz, E. Ferrari, Y. Tao, E. Bertino, and G. Trajcevski. IEEE Computer Society, 2014, pp. 592–603. DOI: 10.1109/ICDE.2014.6816684.
- [7] A. Ailamaki, D. J. DeWitt, and M. D. Hill. “Data page layouts for relational databases on deep memory hierarchies.” In: *The VLDB Journal* 11.3 (Nov. 2002), pp. 198–215. ISSN: 1066-8888. DOI: 10.1007/s00778-002-0074-9.